

Data Preprocessing and Model Evaluation for Flight Price Prediction

Overview

In this script, we prepare the data for a flight price prediction model, using a series of steps including data preprocessing, feature engineering, feature selection using a genetic algorithm, and model evaluation. The model ultimately uses Gradient Boosting Regression to predict flight prices.

Steps Overview

1. Data Preprocessing

- Handle missing values
- Convert categorical features into numerical representations
- Drop unnecessary columns
- Scale numerical features

2. Feature Selection

- Implementing Genetic Algorithm for feature selection

3. Model Training

- Gradient Boosting Regressor

4. Model Evaluation

- Calculate MSE and R2 for model evaluation

1. Data Preprocessing

1.1 Importing Required Libraries

We start by importing the necessary libraries:

```
import pandas as pd
```

```
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV
import random
import warnings
import matplotlib.pyplot as plt
```

1.2 Data Loading and Initial Inspection

After loading the dataset (not shown here), perform the initial inspection:

```
data = pd.read_csv('path_to_file.csv') # Replace with actual path

# Inspect the first few rows of the dataset
print(data.head())

# Inspect the data types and check for missing values
print(data.info())
```

1.3 Handling Ordinal Categorical Features

The **Total_Stops** column is an ordinal categorical feature, and we convert it into numerical form using a dictionary.

```
stops_mapping = {
    'Non-Stop': 0,
    '1 Stop': 1,
    '2 Stops': 2,
    '3 Stops': 3,
    '4 Stops': 4
}
data['Total_Stops'] = data['Total_Stops'].map(stops_mapping)
```

1.4 Dropping Irrelevant Features

We drop the **Duration** column, as it's not relevant to the prediction task.

```
data.drop(columns=['Duration'], inplace=True, errors='ignore')
```

1.5 Handling Datetime Features

Datetime features are converted into numerical format for model compatibility.

```
datetime_cols = data.select_dtypes(include=['datetime64[ns]']).columns
```

```
for col in datetime_cols:
    data[col] = pd.to_numeric(data[col])
```

1.6 Data Splitting

We split the data into features (**X**) and the target variable (**y**), then further split the data into training and test sets.

```
X = data.drop('Price', axis=1)
y = data['Price']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

1.7 Handling Missing Values and Scaling

A pipeline is created to handle missing values and scale numerical features.

```
numeric_cols = X_train.select_dtypes(include=np.number).columns
preprocessor = Pipeline([
    ('imputer', SimpleImputer(strategy='mean')), # Impute missing values
    ('scaler', StandardScaler())               # Scale numerical features
])
```

```
X_train[numeric_cols] = pd.DataFrame(preprocessor.fit_transform(X_train[numeric_cols]),
                                     index=X_train.index, columns=numeric_cols)
X_test[numeric_cols] = pd.DataFrame(preprocessor.transform(X_test[numeric_cols]),
                                   index=X_test.index, columns=numeric_cols)
```

2. Feature Selection Using Genetic Algorithm

The genetic algorithm (GA) is implemented for automatic feature selection. It evaluates subsets of features to identify the best performing ones.

2.1 Genetic Algorithm Functions

```
def evaluate_chromosome(chromosome, X_train, y_train, X_test, y_test):
    selected_features = X_train.columns[chromosome == 1]
    if len(selected_features) == 0:
        return float('inf'), 0 # Penalize empty feature sets

    X_train_subset = X_train[selected_features]
    X_test_subset = X_test[selected_features]

    model = LinearRegression()
    model.fit(X_train_subset, y_train)
    y_pred = model.predict(X_test_subset)
    mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)
    return mse, r2
```

Other GA functions (selection, crossover, mutation) follow the same structure and logic.

2.2 Running the Genetic Algorithm

```
best_features_mask, best_mse, best_r2, all_fitnesses = genetic_algorithm(X_train, y_train,
X_test, y_test)
selected_features = X_train.columns[best_features_mask == 1]
```

3. Model Training Using Gradient Boosting Regressor

3.1 Hyperparameter Tuning Using GridSearchCV

```
param_grid = {
    'n_estimators': [100, 200, 300],
    'learning_rate': [0.01, 0.05, 0.1],
    'max_depth': [3, 4, 5],
    'random_state': [42]
}

gbr = GradientBoostingRegressor()
grid_search = GridSearchCV(gbr, param_grid, cv=3, scoring='neg_mean_squared_error',
verbose=1)

grid_search.fit(X_train_selected, y_train)
best_gbr = grid_search.best_estimator_
```

3.2 Making Predictions

```
y_pred_gbr = best_gbr.predict(X_test_selected)
```

3.3 Model Evaluation

```
mse_gbr = mean_squared_error(y_test, y_pred_gbr)
r2_gbr = r2_score(y_test, y_pred_gbr)
```

```
print(f"Mean Squared Error (MSE): {mse_gbr:.2f}")
print(f"R-squared (R2): {r2_gbr:.2f}")
```

4. Model Evaluation Visualization

```
# Plot the fitness over generations
plt.figure(figsize=(10, 6))
plt.plot([min([f[0] for f in gen_fitnesses]) for gen_fitnesses in all_fitnesses])
plt.xlabel("Generation")
plt.ylabel("Minimum MSE")
plt.title("Genetic Algorithm Convergence")
plt.grid(True)
plt.show()
```

Summary of Steps

1. **Data Preprocessing:** The dataset was cleaned, missing values imputed, and categorical features encoded. Unnecessary features were dropped, and datetime features were converted into numerical values.
2. **Feature Selection:** A genetic algorithm was used to select the best subset of features based on the model's performance, using the Mean Squared Error (MSE) and R-squared (R2) metrics.
3. **Model Training and Hyperparameter Tuning:** A Gradient Boosting Regressor model was trained on the selected features with hyperparameter tuning using GridSearchCV.
4. **Model Evaluation:** The final model was evaluated using MSE and R2, and its performance was visualized through a plot showing the convergence of the genetic algorithm.

Conclusion

The model has been successfully trained, and the features have been optimally selected using a genetic algorithm. The Gradient Boosting Regressor performed well, providing accurate predictions for flight prices based on the selected features.

RESULT



